

Xv6 System Calls: Implementation Guide

Xv6 Kernel Development Team

May 2025

Contents

1	Introduction	3
2	System Call: <code>getinterruptcount</code>	3
2.1	Overview	3
2.2	Implementation Steps	3
2.3	Kernel Code	3
2.4	System Call Registration	4
2.5	User Program Example	4
3	System Call: <code>getdiskstats</code>	4
3.1	Overview	4
3.2	Implementation Steps	4
3.3	Kernel Code	5
3.4	System Call Registration	5
3.5	User Program Example	6
4	System Call: <code>getprocinfo</code>	6
4.1	Overview	6
4.2	Implementation Steps	6
4.3	Kernel Code	6
4.4	System Call Registration	7
4.5	User Program Example	8
5	System Call: <code>getsysinfo</code>	8
5.1	Overview	8
5.2	Implementation Steps	8
5.3	Kernel Code	8
5.4	System Call Registration	9
5.5	User Program Example	9
6	System Call: <code>top</code>	10
6.1	Overview	10
6.2	Implementation Steps	10
6.3	Kernel Code	10
6.4	System Call Registration	11
6.5	User Program Example	12
7	Makefile Modifications	12

1 Introduction

This document provides a detailed guide for implementing five system calls in the Xv6 operating system: `getinterruptcount`, `getdiskstats`, `getprocinfo`, `getsysinfo`, and `top`. Each system call enhances Xv6's functionality by providing user-space programs access to kernel-level information, such as interrupt counts, disk statistics, process details, system resource usage, and top CPU-consuming processes. The implementations include kernel modifications, system call handlers, and user-space interfaces.

2 System Call: `getinterruptcount`

2.1 Overview

The `getinterruptcount` system call returns the total number of hardware interrupts handled by the Xv6 kernel since boot, useful for monitoring kernel activity or debugging.

2.2 Implementation Steps

1. **Declare Global Interrupt Counter:** Add a global variable in `kernel/trap.c` to track interrupts.
2. **Increment Counter in Interrupt Handler:** Modify `kerneltrap` to increment the counter for each interrupt.
3. **Create Kernel Function:** Define `get_interrupt_count` to return the counter value.
4. **Add System Call Interface:** Implement the system call handler and register it.
5. **User-Space Integration:** Add user-space declarations and stubs.

2.3 Kernel Code

Listing 1: `kernel/trap.c`: Global Counter and Handler

```
1 #include "types.h"
2 #include "spinlock.h"
3
4 uint64 interrupt_count = 0;
5
6 void kerneltrap() {
7     // Existing code ...
8     if ((which_dev = devintr()) == 0) {
9         panic("kerneltrap");
10    }
11    interrupt_count++; // Increment for each recognized interrupt
12    // Existing code ...
13 }
14
15 uint64 get_interrupt_count(void) {
16     return interrupt_count;
17 }
```

Listing 2: `kernel/sysproc.c`: System Call Handler

```
1 uint64 sys_getinterruptcount(void) {
2     return get_interrupt_count();
3 }
```

Listing 3: kernel/defs.h: Function Prototype

```
1 uint64 get_interrupt_count(void);
```

2.4 System Call Registration

Listing 4: kernel/syscall.h: Syscall Number

```
1 #define SYS_getinterruptcount 22
```

Listing 5: kernel/syscall.c: Syscall Table

```
1 extern uint64 sys_getinterruptcount(void);
2 [SYS_getinterruptcount] sys_getinterruptcount,
```

Listing 6: user/usys.pl: User Stub

```
1 entry("getinterruptcount")
```

Listing 7: user/user.h: User Declaration

```
1 uint64 getinterruptcount(void);
```

2.5 User Program Example

Listing 8: user/interrupt.c: Test Program

```
1 #include "user.h"
2 #include "stdio.h"
3 int main() {
4     uint64 count = getinterruptcount();
5     printf("Total interrupts: %llu\n", count);
6     return 0;
7 }
```

3 System Call: getdiskstats

3.1 Overview

The `getdiskstats` system call retrieves disk I/O statistics, including read/write counts and bytes, from the `virtio` disk driver.

3.2 Implementation Steps

1. **Define Disk Stats Structure:** Create `struct diskstats` in `kernel/virtio.h`.
2. **Add Global Counters:** Track statistics in `kernel/virtio_disk.c`. **Implement Kernel Function:** `Copystatstouserspace` in `kernel/sysproc.c`.
3. **Register System Call:** Add syscall number and handler.
4. **User-Space Wrapper:** Provide user-space declarations and test program.

3.3 Kernel Code

Listing 9: kernel/virtio.h: Disk Stats Structure

```
1 struct diskstats {
2     uint64 read_count;
3     uint64 write_count;
4     uint64 read_bytes;
5     uint64 write_bytes;
6 };
```

Listing 10: kernel/virtio_{disk}.c : *GlobalCounters*

```
1 #include "virtio.h"
2 struct diskstats disk_stats = {0, 0, 0, 0};
3
4 // In read completion:
5 disk_stats.read_count++;
6 disk_stats.read_bytes += number_of_bytes_read;
7
8 // In write completion:
9 disk_stats.write_count++;
10 disk_stats.write_bytes += number_of_bytes_written;
```

Listing 11: kernel/sysproc.c: System Call Handler

```
1 #include "virtio.h"
2 #include "defs.h"
3 #include "proc.h"
4
5 uint64 sys_getdiskstats(void) {
6     struct diskstats stats = disk_stats;
7     struct diskstats *user_ptr;
8     argaddr(0, (uint64*)&user_ptr);
9     if (copyout(myproc()->pagetable, (uint64)user_ptr, (char *)&stats,
10         sizeof(stats)) < 0)
11         return -1;
12     return 0;
13 }
```

3.4 System Call Registration

Listing 12: kernel/syscall.h: Syscall Number

```
1 #define SYS_getdiskstats 23
```

Listing 13: kernel/syscall.c: Syscall Table

```
1 extern uint64 sys_getdiskstats(void);
2 [SYS_getdiskstats] sys_getdiskstats,
```

Listing 14: user/usys.pl: User Stub

```
1 entry("getdiskstats")
```

Listing 15: user/user.h: User Declaration

```
1 struct diskstats {
2     uint64 read_count;
```

```

3   uint64 write_count;
4   uint64 read_bytes;
5   uint64 write_bytes;
6 };
7 int getdiskstats(struct diskstats *stats);

```

3.5 User Program Example

Listing 16: user/diskinfo.c: Test Program

```

1 #include "user.h"
2 #include "diskstats.h"
3 #include "stdio.h"
4 int main(void) {
5     struct diskstats stats;
6     if (getdiskstats(&stats) == 0) {
7         printf("Disk reads: %d\n", stats.read_count);
8         printf("Disk writes: %d\n", stats.write_count);
9         printf("Bytes read: %d\n", stats.read_bytes);
10        printf("Bytes written: %d\n", stats.write_bytes);
11    } else {
12        printf("Failed to get disk stats\n");
13    }
14    return 0;
15 }

```

4 System Call: getprocinfo

4.1 Overview

The `getprocinfo` system call provides information about a specific process, given its PID, including state, parent PID, memory size, and name.

4.2 Implementation Steps

1. **Define procinfo Structure:** Create `struct procinfo` in `kernel/sysinfo.h`.
2. **Kernel Function:** Implement `getprocinfo` in `kernel/proc.c`.
3. **System Call Wrapper:** Add handler in `kernel/sysproc.c`.
4. **Register System Call:** Add syscall number and mappings.
5. **User-Space Test:** Provide user-space program to test the call.

4.3 Kernel Code

Listing 17: kernel/sysinfo.h: procinfo Structure

```

1 #ifndef SYSINFO_H
2 #define SYSINFO_H
3 struct procinfo {
4     int pid;
5     int state;
6     int ppid;
7     uint64 sz;
8     char name[16];
9 };

```

```
10 #endif
```

Listing 18: kernel/proc.c: Kernel Function

```
1 #include "sysinfo.h"
2 uint64 getprocinfo(int pid, struct procinfo *info) {
3     struct proc *p;
4     for (p = proc; p < &proc[NPROC]; p++) {
5         if (p->state != UNUSED && p->pid == pid) {
6             info->pid = p->pid;
7             info->state = p->state;
8             info->ppid = p->parent ? p->parent->pid : 0;
9             info->sz = p->sz;
10            safestrcpy(info->name, p->name, sizeof(p->name));
11            return 0;
12        }
13    }
14    return -1;
15 }
```

Listing 19: kernel/sysproc.c: System Call Handler

```
1 #include "sysinfo.h"
2 uint64 sys_getprocinfo(void) {
3     int pid;
4     uint64 user_addr;
5     struct procinfo info;
6     argint(0, &pid);
7     argaddr(1, &user_addr);
8     if (getprocinfo(pid, &info) < 0)
9         return -1;
10    if (copyout(myproc()->pagetable, user_addr, (char *)&info, sizeof(
11        info)) < 0)
12        return -1;
13    return 0;
14 }
```

4.4 System Call Registration

Listing 20: kernel/syscall.h: Syscall Number

```
1 #define SYS_getprocinfo 24
```

Listing 21: kernel/syscall.c: Syscall Table

```
1 extern uint64 sys_getprocinfo(void);
2 [SYS_getprocinfo] sys_getprocinfo,
```

Listing 22: user/usys.pl: User Stub

```
1 entry("getprocinfo")
```

Listing 23: user/user.h: User Declaration

```
1 #include "sysinfo.h"
2 int getprocinfo(int pid, struct procinfo *info);
```

4.5 User Program Example

Listing 24: user/sysinfo.c: Test Program

```
1 #include "kernel/types.h"
2 #include "user/user.h"
3 int main(int argc, char *argv[]) {
4     struct procinfo info;
5     int pid = 1;
6     if (getprocinfo(pid, &info) == 0) {
7         printf("PID: %d, State: %d, PPID: %d, Size: %d, Name: %s\n",
8             info.pid, info.state, info.ppid, info.sz, info.name);
9     } else {
10        printf("Process %d not found\n", pid);
11    }
12    exit(0);
13 }
```

5 System Call: getsysinfo

5.1 Overview

The `getsysinfo` system call returns system-wide information, including the number of processes, runnable processes, and free memory in bytes.

5.2 Implementation Steps

1. **Define sysinfo Structure:** Create `struct sysinfo` in `kernel/sysinfo.h`.
2. **Count Free Pages:** Add `freepages` in `kernel/kalloc.c`.
3. **Implement System Call:** Add handler in `kernel/sysproc.c`.
4. **Register System Call:** Add syscall number and mappings.
5. **User Program:** Create a test program.

5.3 Kernel Code

Listing 25: kernel/sysinfo.h: sysinfo Structure

```
1 #ifndef SYSINFO_H
2 #define SYSINFO_H
3 struct sysinfo {
4     int nprocs;
5     int nrunnable;
6     int freemem;
7 };
8 #endif
```

Listing 26: kernel/kalloc.c: Free Pages Counter

```
1 int freepages(void) {
2     struct run *r;
3     int count = 0;
4     acquire(&kmem.lock);
5     for (r = kmem.freelist; r; r = r->next)
6         count++;
7     release(&kmem.lock);
8 }
```



```

8     return count;
9 }

```

Listing 27: kernel/sysproc.c: System Call Handler

```

1 #include "sysinfo.h"
2 uint64 sys_getsysinfo(void) {
3     struct sysinfo info;
4     struct proc *p;
5     uint64 addr;
6     info.nprocs = 0;
7     info.nrunnable = 0;
8     info.freemem = freepages() * PGSIZE;
9     for (p = proc; p < &proc[NPROC]; p++) {
10         if (p->state != UNUSED)
11             info.nprocs++;
12         if (p->state == RUNNABLE)
13             info.nrunnable++;
14     }
15     argaddr(0, &addr);
16     if (copyout(myproc()->pagetable, addr, (char *)&info, sizeof(info)) <
17         0)
18         return -1;
19     return 0;
20 }

```

5.4 System Call Registration

Listing 28: kernel/syscall.h: Syscall Number

```

1 #define SYS_getsysinfo 25

```

Listing 29: kernel/syscall.c: Syscall Table

```

1 extern uint64 sys_getsysinfo(void);
2 [SYS_getsysinfo] sys_getsysinfo,

```

Listing 30: user/usys.pl: User Stub

```

1 entry("getsysinfo")

```

Listing 31: user/user.h: User Declaration

```

1 #include "../kernel/sysinfo.h"
2 int getsysinfo(struct sysinfo *info);

```

5.5 User Program Example

Listing 32: user/sysinfo.c: Test Program

```

1 #include "user.h"
2 #include "../kernel/sysinfo.h"
3 int main(void) {
4     struct sysinfo info;
5     if (getsysinfo(&info) < 0) {
6         printf("sysinfo call failed\n");
7         exit(1);
8     }

```

```

9   printf("Processes: %d\n", info.nprocs);
10  printf("Runnable: %d\n", info.nrunnable);
11  printf("Free mem: %d bytes\n", info.freemem);
12  exit(0);
13 }

```

6 System Call: top

6.1 Overview

The `top` system call returns information about the top CPU-consuming processes, including PID, CPU time, and name, sorted by CPU usage.

6.2 Implementation Steps

1. **Define top_proc Structure:** Add `struct top_proc` and `cputime` field in `kernel/proc.h`.
2. **Track CPU Time:** Modify the scheduler in `kernel/proc.c`.
3. **Implement gettop Function:** Add sorting logic in `kernel/proc.c`.
4. **System Call Handler:** Implement `sys_top` in `kernel/sysproc.c`.
5. **Register System Call:** Add syscall number and mappings.
6. **User Program:** Create a test program.

6.3 Kernel Code

Listing 33: `kernel/proc.h`: `top_proc` Structure

```

1  struct top_proc {
2      int pid;
3      uint cputime;
4      char name[16];
5  };
6  struct proc {
7      // ...
8      uint cputime; // Total CPU time used
9      // ...
10 };
11 int gettop(struct top_proc *tops, int n);

```

Listing 34: `kernel/proc.c`: Scheduler and `gettop`

```

1  void scheduler(void) {
2      struct proc *p;
3      for(;;) {
4          sti();
5          acquire(&ptable.lock);
6          for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
7              if(p->state != RUNNABLE)
8                  continue;
9              proc = p;
10             switchvm(p);
11             p->state = RUNNING;
12             switch(&cpu->scheduler, proc->context);
13             switchkvm();

```

```

14     p->cputime++; // Increment CPU time
15     proc = 0;
16 }
17 release(&ptable.lock);
18 }
19 }
20
21 int gettop(struct top_proc *tops, int n) {
22     struct proc *p;
23     int i, j;
24     acquire(&ptable.lock);
25     for (i = 0; i < n; i++) {
26         tops[i].pid = 0;
27         tops[i].cputime = 0;
28         memset(tops[i].name, 0, sizeof(tops[i].name));
29     }
30     for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {
31         if(p->state == UNUSED)
32             continue;
33         for(i = 0; i < n; i++) {
34             if(p->cputime > tops[i].cputime) {
35                 for(j = n - 1; j > i; j--) {
36                     tops[j] = tops[j-1];
37                 }
38                 tops[i].pid = p->pid;
39                 tops[i].cputime = p->cputime;
40                 safestrcpy(tops[i].name, p->name, sizeof(tops[i].name));
41                 break;
42             }
43         }
44     }
45     release(&ptable.lock);
46     return n;
47 }

```

Listing 35: kernel/sysproc.c: System Call Handler

```

1 int sys_top(void) {
2     struct top_proc *user_tops;
3     int n;
4     if (argptr(0, (void*)&user_tops, sizeof(struct top_proc) * 64) < 0 ||
5         argint(1, &n) < 0) {
6         return -1;
7     }
8     struct top_proc kernel_tops[n];
9     int count = gettop(kernel_tops, n);
10    if (copyout(myproc()->pagetable, (uint64)user_tops, (char*)
11        kernel_tops, sizeof(struct top_proc) * count) < 0) {
12        return -1;
13    }
14    return count;
15 }

```

6.4 System Call Registration

Listing 36: kernel/syscall.h: Syscall Number

```

1 #define SYS_top 26

```

Listing 37: kernel/syscall.c: Syscall Table

```
1 extern uint64 sys_top(void);
2 [SYS_top] sys_top,
```

Listing 38: user/usys.pl: User Stub

```
1 entry("top")
```

Listing 39: user/user.h: User Declaration

```
1 struct top_proc {
2     int pid;
3     unsigned int cputime;
4     char name[16];
5 };
6 int top(struct top_proc *tops, int n);
```

6.5 User Program Example

Listing 40: user/top.c: Test Program

```
1 #include "user.h"
2 #include "fcntl.h"
3 #include "types.h"
4 #include "stat.h"
5 int main(int argc, char *argv[]) {
6     struct top_proc tops[10];
7     int n = 10;
8     int count = top(tops, n);
9     if (count < 0) {
10         printf("top syscall failed\n");
11         exit(1);
12     }
13     printf("Top %d processes by CPU time:\n", count);
14     for (int i = 0; i < count; i++) {
15         printf("PID: %d, CPU Time: %d, Name: %s\n", tops[i].pid, tops[i].
16             cputime, tops[i].name);
17     }
18     exit(0);
19 }
```

7 Makefile Modifications

To compile the user programs, update **Makefile** to include the new executables:

Listing 41: Makefile: User Programs

```
1 UPROGS=\
2     _cat\
3     _echo\
4     _forktest\
5     _interrupt\
6     _diskinfo\
7     _sysinfo\
8     _top\
```

8 Conclusion

This document detailed the implementation of five system calls in Xv6: `getinterruptcount`, `getdiskstats`, `getprocinfo`, `getsysinfo`, and `top`. Each system call was implemented with kernel modifications, system call handlers, and user-space test programs, enabling robust system monitoring and introspection.